

考題 for 機器學習專家

- Quiz time: 120 minutes
 - we suggested you spend 40 minutes in part 1 and 80 minutes in part 2

Part 1: Maths and Common Senses (50%)

Quiz 1-1: Back Propagation (20%)

- Given an neural network with one hidden layer, and the loss function is cross entropy loss
 - $a = W^{(1)}x + b^{(1)}$
 - $h = \sigma(a)$
 - $\theta = W^{(2)}h + b^{(2)}$
 - $\hat{y} = softmax(\theta)$
 - $L = CE(y, \hat{y})$
 - $x \in \mathbb{R}^n, W^{(1)} \in \mathbb{R}^{d \times n}, b^{(1)} \in \mathbb{R}^d, h \in \mathbb{R}^d, W^{(2)} \in \mathbb{R}^{c \times d}, b^{(2)} \in \mathbb{R}^c, \theta \in \mathbb{R}^c$
- Calculate $\frac{\partial}{\partial \theta} L$ and $\frac{\partial}{\partial W^{(2)}} L$

Note: you can write answers on a paper and take a photo of the paper

Quiz 1-2: Common Senses (30%)

Note: you can answer in Chinese or English.

Note: your answers can be as simplest as possible

- When will we use F1-score instead of precision(accuracy)?
- Why don't we use binary classification function as the activation function in neural networks?
 - Note: binary classification function $f(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$
- What is the bias and variance of a machine learning algorithm?
- When training a single tree in random forest, we don't prune the tree, why?
- What is one-hot encoding?
- How to prevent overfitting in neural networks? (write down anything you know)

Part 2: Python Skills (50%)

Quiz 2-1: word counts (10%)

Download this [English article](https://goo.gl/BDB6bE) from `https://goo.gl/BDB6bE` first.

We want to know the probabilities of each bigram and trigram.

Please write a function `ngram_probs` to calculate these counts.

Note: please convert text to lower cases before calculating.

Note: please keep punctuations in the text.

Note: in all quizzes, you can use any Python packages on PyPI.

```
1 def ngram_probs(filename='raw_sentences.txt'):
2     return bigram_probs, trigram_probs
```

```

3
4 cnt2, cnt3 = ngram_probs()
5 print(cnt2[('we', 'are')])

```

Note: an n-gram is a contiguous sequence of n items from a given sequence of text or speech. For example, ('we',) is an unigram, ('we', 'are') is a bigram, and ('we', 'are', 'here') is a trigram.

Quiz 2-2: next word probabilities (10%)

Write a function `prob3` to calculate the probabilities of next word of a bigram.

The probability of 'family' after bigram ('we', 'are') is the probability ('we', 'are', 'family') divided by the probability ('we', 'are').

$$P(w_3|w_1, w_2) = \frac{P(w_1, w_2, w_3)}{P(w_1, w_2)}$$

$$P(\text{family}|\text{we are}) = \frac{P(\text{we are family})}{P(\text{we are})}$$

Note: In this quiz, you are requested to return **log probabilities**.

```

1 def prob3(bigram, cnt2=cnt2, cnt3=cnt3):
2     return prob
3
4 p = prob3(('we', 'are'))
5 print(p['family'])

```

Quiz 2-3: predicting the next word (10%)

Write a function `predict_max` to complete the sentence by finding the max likelihood word.

For example, if the starting bigram is ('we', 'are'), the max likelihood word is 'going'.

Given the next word 'going', current bigram becomes ('are', 'going'), and then we can predict the next word. The predicting process will end when the next word is '.' or length of the sentence is larger than 15.

```

1 def predict_max(starting, cnt2=cnt2, cnt3=cnt3):
2     return list_of_words
3
4 sent = predict_max(('we', 'are'))
5 assert sent[-1] == '.' or len(sent) <= 15
6 print(' '.join(sent))

```

Quiz 2-4: beam search (20%)

Write a function `predict_beam` to complete the sentence by **beam search**

https://en.wikipedia.org/wiki/Beam_search. "Complete" means find all sentences which end in "." and have `sent_length` tokens (including ".").

```
1 def predict_beam(bigram, beam_size=4, sent_length=10, cnt2=cnt2, cnt3=cnt
  3):
2     return list_of_sentence
3
4 for sent in predict_beam(('we', 'are')):
5     assert sent[-1] == '.' or len(sent) < 10
6     print(' '.join(sent))
```